

Homework 2

Aiden Kenny

STAT GR5205: Linear Regression Models

Columbia University

October 5, 2020

Question 1 *Collaborators:* None

Supposed for $\mathbf{x}, \mathbf{y}, \boldsymbol{\epsilon}, \mathbf{1} \in \mathbb{R}^n$, where $\mathbf{x}, \mathbf{1}$ are *fixed* vectors and $\mathbf{y}, \boldsymbol{\epsilon}$ are *random* vectors, the simple linear regression model

$$\mathbf{y} = \beta_0 \mathbf{1} + \beta_1 \mathbf{x} + \boldsymbol{\epsilon}$$

holds, with $\mathbb{E}[\boldsymbol{\epsilon}] = \mathbf{0}$ and $\text{Var}[\boldsymbol{\epsilon}] = \sigma^2 \mathbf{I}$. The least-squares estimators are given by

$$\hat{\beta}_1 = \frac{(\mathbf{x} - \bar{x}\mathbf{1})^T(\mathbf{y} - \bar{y}\mathbf{1})}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2}, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}, \quad \text{and} \quad \hat{\sigma}^2 = \frac{1}{n-2} \left\| \mathbf{y} - \hat{\beta}_0 \mathbf{1} - \hat{\beta}_1 \mathbf{x} \right\|^2.$$

- (a) We first determine several properties of \mathbf{y} (a random vector) and \bar{y} (a random variable). For \mathbf{y} , we have

$$\begin{aligned} \mathbb{E}[\mathbf{y}] &= \mathbb{E}[\beta_0 \mathbf{1} + \beta_1 \mathbf{x} + \boldsymbol{\epsilon}] = \beta_0 \mathbf{1} + \beta_1 \mathbf{x} + \mathbb{E}[\boldsymbol{\epsilon}] = \beta_0 \mathbf{1} + \beta_1 \mathbf{x}, \\ \text{Var}[\mathbf{y}] &= \text{Var}[\beta_0 \mathbf{1} + \beta_1 \mathbf{x} + \boldsymbol{\epsilon}] = \mathbf{0} + \mathbf{0} + \text{Var}[\boldsymbol{\epsilon}] = \sigma^2 \mathbf{I}. \end{aligned}$$

That is, for each y_i , we have $\mathbb{E}[y_i] = \beta_0 + \beta_1 x_i$ and $\text{Var}[y_i] = \sigma^2$. We also have $\text{Cov}[y_i, y_j] = 0$ for all $i \neq j$. For \bar{y} , we have

$$\begin{aligned} \mathbb{E}[\bar{y}] &= \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n y_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[y_i] = \frac{1}{n} \sum_{i=1}^n (\beta_0 + \beta_1 x_i) = \beta_0 + \beta_1 \bar{x}, \\ \text{Var}[\bar{y}] &= \text{Var}\left[\frac{1}{n} \sum_{i=1}^n y_i\right] = \frac{1}{n^2} \sum_{i=1}^n \text{Var}[y_i] = \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{\sigma^2}{n}. \end{aligned}$$

Expanding out $\hat{\beta}_1$ gives us

$$\hat{\beta}_1 = \frac{(\mathbf{x} - \bar{x}\mathbf{1})^T(\mathbf{y} - \bar{y}\mathbf{1})}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} = \frac{1}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}).$$

For notational ease, we are going to multiply both sides of this estimate by $\|\mathbf{x} - \bar{x}\mathbf{1}\|^2$, since it is just a constant. Taking the expected value of $\|\mathbf{x} - \bar{x}\mathbf{1}\|^2 \cdot \hat{\beta}_1$ yields

$$\begin{aligned} \mathbb{E}\left[\|\mathbf{x} - \bar{x}\mathbf{1}\|^2 \cdot \hat{\beta}_1\right] &= \|\mathbf{x} - \bar{x}\mathbf{1}\|^2 \cdot \mathbb{E}[\hat{\beta}_1] = \mathbb{E}\left[\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})\right] \\ &= \sum_{i=1}^n \mathbb{E}[(x_i - \bar{x})(y_i - \bar{y})] = \sum_{i=1}^n (x_i - \bar{x})(\mathbb{E}[y_i] - \mathbb{E}[\bar{y}]) \\ &= \sum_{i=1}^n (x_i - \bar{x})(\beta_0 + \beta_1 x_i - \beta_0 - \beta_1 \bar{x}) = \beta_1 \sum_{i=1}^n (x_i - \bar{x})^2 = \beta_1 \cdot \|\mathbf{x} - \bar{x}\mathbf{1}\|^2. \end{aligned}$$

Dividing both sides of the equation shows that $\mathbb{E}[\hat{\beta}_1] = \beta_1$. Next, taking the expected value of $\hat{\beta}_0$ gives us

$$\mathbb{E}[\hat{\beta}_0] = \mathbb{E}[\bar{y} - \hat{\beta}_1 \bar{x}] = \mathbb{E}[\bar{y}] - \bar{x} \mathbb{E}[\hat{\beta}_1] = \beta_0 + \beta_1 \bar{x} - \beta_1 \bar{x} = \beta_0.$$

Finally, taking the expected value of $\hat{\sigma}^2$ leads to

(b) Looking at the expanded equation for $\hat{\beta}_1$, we have

$$\begin{aligned}\|\mathbf{x} - \bar{x}\mathbf{1}\|^2 \cdot \hat{\beta}_1 &= \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) = \sum_{i=1}^n (x_i - \bar{x})y_i - \bar{y} \sum_{i=1}^n (x_i - \bar{x}) \\ &= \sum_{i=1}^n (x_i - \bar{x})y_i - \bar{y}(n\bar{x} - n\bar{x}) = \sum_{i=1}^n (x_i - \bar{x})y_i.\end{aligned}$$

That is, we are able to remove the \bar{y} from the summation entirely. By taking the variance of $\|\mathbf{x} - \bar{x}\mathbf{1}\|^2 \cdot \hat{\beta}_1$, we have

$$\begin{aligned}\text{Var}\left[\|\mathbf{x} - \bar{x}\mathbf{1}\|^2 \cdot \hat{\beta}_1\right] &= \|\mathbf{x} - \bar{x}\mathbf{1}\|^4 \cdot \text{Var}[\hat{\beta}_1] = \text{Var}\left[\sum_{i=1}^n (x_i - \bar{x})y_i\right] \\ &= \sum_{i=1}^n \text{Var}[(x_i - \bar{x})y_i] = \sum_{i=1}^n (x_i - \bar{x})^2 \text{Var}[y_i] = \sum_{i=1}^n (x_i - \bar{x})^2 \sigma^2 = \sigma^2 \cdot \|\mathbf{x} - \bar{x}\mathbf{1}\|^2,\end{aligned}$$

and dividing both sides by $\|\mathbf{x} - \bar{x}\mathbf{1}\|^2$ shows that $\text{Var}[\hat{\beta}_1] = \sigma^2 / \|\mathbf{x} - \bar{x}\mathbf{1}\|^2$. Similarly, taking the variance of $\hat{\beta}_0$ gives us

$$\text{Var}[\hat{\beta}_0] = \text{Var}[\bar{y} - \hat{\beta}_1 \bar{x}] = \text{Var}[\bar{y}] + \bar{x}^2 \text{Var}[\hat{\beta}_1] = \frac{\sigma^2}{n} + \bar{x}^2 \frac{\sigma^2}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{\|\mathbf{x} - \bar{x}\mathbf{1}\|^2} \right).$$

Question 2 *Collaborators:* None

Letting $\boldsymbol{\beta} = (\beta_0, \beta_1)^T$ and $\mathbf{X} = [\mathbf{1} \ \mathbf{x}] \in \mathbb{R}^{n \times 2}$, the simple linear regression model is given by $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$. The MSE is then given by $Q = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2$.

(a) We first expand the MSE to get

$$Q = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2 = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) = \mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\beta}^T \mathbf{X}^T \mathbf{X}\boldsymbol{\beta}.$$

By differentiating Q with respect to $\boldsymbol{\beta}$ and setting it equal to $\mathbf{0}$, we have

$$\begin{aligned}\frac{\partial Q}{\partial \boldsymbol{\beta}} &= \frac{\partial}{\partial \boldsymbol{\beta}} (\mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\beta}^T \mathbf{X}^T \mathbf{X}\boldsymbol{\beta}) \\ &= \frac{\partial}{\partial \boldsymbol{\beta}} \mathbf{y}^T \mathbf{y} - 2 \frac{\partial}{\partial \boldsymbol{\beta}} \mathbf{y}^T \mathbf{X}\boldsymbol{\beta} + \frac{\partial}{\partial \boldsymbol{\beta}} \boldsymbol{\beta}^T \mathbf{X}^T \mathbf{X}\boldsymbol{\beta} = -2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X}\boldsymbol{\beta} \stackrel{\text{set}}{=} \mathbf{0},\end{aligned}$$

and solving for $\boldsymbol{\beta}$ gives us $\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$.

(b) Now suppose that we have p different predictors, each with n observed values that are not all identical. Let $\mathbf{x}_j \in \mathbb{R}^n$ be the vector containing the observations for the j th predictor. Define the matrix $\mathbf{X} = [\mathbf{1} \ \mathbf{x}_1 \ \cdots \ \mathbf{x}_p] \in \mathbb{R}^{n \times (p+1)}$ be the matrix whos j th column is \mathbf{x}_{j-1} (and first column is $\mathbf{1}$). In addition, let $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^T \in \mathbb{R}^{p+1}$ be the vector containing $p+1$ scalars. The multiple linear regression model is given by $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$; that is, it has the same form as the simple linear regression model, and we can see that simple linear regression is when $p = 1$. Because of this, the estimated coefficients $\hat{\boldsymbol{\beta}}$ take the same form as before: $\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$.

(c) The fitted values $\hat{\mathbf{y}}$ are given by

$$\hat{\mathbf{y}} = \mathbf{X}\hat{\boldsymbol{\beta}} = \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

(d) Similar to simple linear regression, the “normal equation” for the multivariate regression setting is found during the derivation of $\hat{\boldsymbol{\beta}}$, and is given by

$$\mathbf{X}^T (\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) = \mathbf{0}.$$

Letting $\mathbf{e} = \mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}$ be the vector of observed residuals, we have $\mathbf{X}^T \mathbf{e} = \mathbf{0}$, meaning $\mathbf{1}^T \mathbf{e} = 0$ and $\mathbf{x}_j^T \mathbf{e} = 0$ for all j . In other words, the residuals sum to zero and the residuals weighted by each predictor sum to 0. Because of this result, we also know that the residuals weighted by the fitted values sum to zero, i.e.

$$\hat{\mathbf{y}}^T \mathbf{e} = (\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y})^T \mathbf{e} = \mathbf{y}^T \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{e} = \mathbf{y}^T \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{0} = \mathbf{0}.$$

The other key results are that $\hat{\mathbf{y}}^T \mathbf{1} = \mathbf{y}^T \mathbf{1}$ and $\bar{y} = \bar{\mathbf{x}}^T \hat{\boldsymbol{\beta}}$, where $\bar{\mathbf{x}} = (1, \bar{x}_1, \dots, \bar{x}_p)^T$ is the vector whos j th entry is the mean of the $(j - 1)$ th variable.

(e) We have

$$\begin{aligned} \|\hat{\mathbf{y}} - \bar{y}\mathbf{1}\|^2 &= (\hat{\mathbf{y}} - \bar{y}\mathbf{1})^T (\hat{\mathbf{y}} - \bar{y}\mathbf{1}) \\ &= \hat{\mathbf{y}}^T \hat{\mathbf{y}} - \bar{y} \hat{\mathbf{y}}^T \mathbf{1} - \bar{y} \hat{\mathbf{y}}^T \mathbf{1} + \bar{y}^2 \mathbf{1}^T \mathbf{1} \\ &= \hat{\mathbf{y}}^T \hat{\mathbf{y}} - \bar{y} \mathbf{y}^T \mathbf{1} - \bar{y} \hat{\mathbf{y}}^T \mathbf{1} + \bar{y}^2 \mathbf{1}^T \mathbf{1} && \text{(since } \hat{\mathbf{y}}^T \mathbf{1} = \mathbf{y}^T \mathbf{1}) \\ &= \hat{\mathbf{y}}^T (\mathbf{y} - \mathbf{e}) - \bar{y} \mathbf{y}^T \mathbf{1} - \bar{y} \hat{\mathbf{y}}^T \mathbf{1} + \bar{y}^2 \mathbf{1}^T \mathbf{1} && \text{(since } \hat{\mathbf{y}} = \mathbf{y} - \mathbf{e}) \\ &= \hat{\mathbf{y}}^T \mathbf{y} - \hat{\mathbf{y}}^T \mathbf{e} - \bar{y} \mathbf{y}^T \mathbf{1} - \bar{y} \hat{\mathbf{y}}^T \mathbf{1} + \bar{y}^2 \mathbf{1}^T \mathbf{1} \\ &= \hat{\mathbf{y}}^T \mathbf{y} - \bar{y} \mathbf{y}^T \mathbf{1} - \bar{y} \hat{\mathbf{y}}^T \mathbf{1} + \bar{y}^2 \mathbf{1}^T \mathbf{1} && \text{(since } \hat{\mathbf{y}}^T \mathbf{e} = \mathbf{0}) \\ &= (\hat{\mathbf{y}} - \bar{y}\mathbf{1})^T (\mathbf{y} - \bar{y}\mathbf{1}). \end{aligned}$$

Essentially, we expanded out the equation and used the key results from part (d) to mainpulate the equation and get the desired result.

Question 3 *Collaborators:* None

- (a) Our hypothesis test is given by $H_0 : \beta_1 \leq 0$ vs. $H_a : \beta_1 > 0$, and we fail to reject H_0 . The analyst claiming that there is no linear association between X and Y is *incorrect*. The data indicates that $\beta_1 \leq 0$, not $\beta_1 = 1$, meaning there could be a negative linear association.
- (b) While this situation is more subtle, the analyst is again incorrect. In the first situation, we are estimating the true value of $\mathbb{E}[Y|X = x_0]$, so we are trying to estimate an underlying feature of the random variable Y . In the second situation, we are trying to predict the average of *random samples* from Y , i.e. we want to estimate $\mathbb{E}[\bar{Y}|X = x_0]$. One reason for this confusion is that the two results are the same. For example, if $Z \sim N(\mu, \sigma^2)$, then $\bar{Z} \sim N(\mu, \sigma^2/n)$, which means $\mathbb{E}[Z] = \mathbb{E}[\bar{Z}] = \mu$. It is important to note that even though the numerical answer is the same, *what* we are doing in each situation is different.
- (c) At the point $X = x_0$, a 95% confidence interval for the mean value of Y_0 is given by

$$\hat{Y}_0 \pm t_{0.975, n-2} \sqrt{\text{MSE} \left(\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum_{i=1}^n (x_i - \bar{x})^2} \right)}.$$

However, if we wanted to determine the range that a *single observation* would fall in, the interval would not be the same for several reasons. We are actually looking for a 95% *prediction interval*, which at some different point $X = x_h \neq x_0$, would be given by

$$\hat{Y}_h \pm t_{0.975, n-2} \sqrt{\text{MSE} \left(1 + \frac{1}{n} + \frac{(x_h - \bar{x})^2}{\sum_{i=1}^n (x_i - \bar{x})^2} \right)}.$$

We also note that the prediction interval at x_h is based around \hat{Y}_h , while the confidence interval at x_o is based around \hat{Y}_o . In addition, a prediction interval is always wider than a confidence interval.

Question 4 *Collaborators:* None

For this question, we are working with the copier dataset, where X denotes the number of copiers serviced and Y denotes the total number of minutes spent on a service call.

- (a) We want to determine whether or not there is a linear relationship between X and Y , i.e. $Y = \beta_0 + \beta_1 X + \epsilon$. To do this, we use R to fit a linear model, which is estimated to be

$$Y = -0.5802 + 15.0352X.$$

When using the `lm()` function, a hypothesis test is automatically carried out with $H_0 : \beta_1 = 0$ and $H_a : \beta_1 \neq 0$ (note that R has other functions that can be used to conduct non-standard tests). The p -value is given as $p < 2.2 \cdot 10^{-16} < 0.05$, so we reject H_0 . The data indicated that there is indeed a linear relationship between X and Y .

- (b) Using the `confint()` function in R, a 95% confidence interval for β_1 is given by

$$\beta_1 \in (14.061010, 16.009486).$$

That is, we believe that there is a 0.95 probability that the true value of β_1 can be no smaller than 14.061010 and no larger than 16.009486.

- (c) The manager claims that the mean service time should not increase by more than 14 minutes for each printer serviced.
1. We can see that the confidence interval for β_1 lies entirely above 14, which indicates that the manager is incorrect.
 2. We can perform a hypothesis test with $H_0 : \beta_1 \leq 14$ and $H_a : \beta_1 > 14$. For this experiment, we have $n = 45$. Our t -statistic is then given by

$$t = \frac{\hat{\beta}_1 - \beta_1}{\hat{\sigma}} = \frac{15.0352 - 14}{0.4831} = 2.142984,$$

and our corresponding p -value is

$$p = \Pr(T \geq 2.142984 | \beta_1 = 14) \approx 0.0189 < 0.05.$$

Since our p -value is smaller than our significance level, we reject H_0 .

Our results are consistent in both cases, and the data indicates that the average service time will increase by more than 14 minutes.

- (d) In this model, β_0 would be the expected service time when no copiers are being serviced, which does not make sense (if no copiers are being serviced, then the customer would have no reason to call in the first place). Even if it did, we have $\hat{\beta}_0 = -0.5802$, a negative service time, which is an impossible value to be observed.